

PEACELAND

LEMONNIER Alexandre

MOLLARD Emmanuel

QUESLATI Mehdi

SIMONIN Victor



TECHNICAL & BUSINESS CONSTRAINTS FOR STATISTICS' DATA STORAGE



Statistics should reliably encompass any piece of data

Consistency




Statistics may come from every corner of Peaceland asynchronously

Partition Tolerance



It is fine if the database is temporarily unavailable to external requests.



Statistics will probably involve heavy data aggregation

Column Oriented

APACHE
HBASE




FAST: FAST ALERT SIGNAL TELLER



Speed

Every second a soul is troubled is a second without peace.

And that's *unacceptable*.



Reactivity

Peacewatchers must be able to fetch data quickly, which involves keeping an AP database.



Coherence

Since peacewatchers may produce reports based on slightly not-up-to-date information, and since those reports may have immediate consequences, the reports that trigger alerts are also prioritized by the system.



SCYLLA

PREVIOUS MISTAKES ON THE FAILED ATTEMPT

Hired data-scientists instead
of data-engineers
(~~Isk~~ they weren't even from Epita)

Used either SQL, or non-
CP, or a centralized
database for their statistics.

Kept using an interpreted
language after the POC.

Bad architecture for the
volume of data, no
streams...

No AP database for day-
to-day activities.

TECHNICAL INFORMATION TO ADD TO THE REPORT

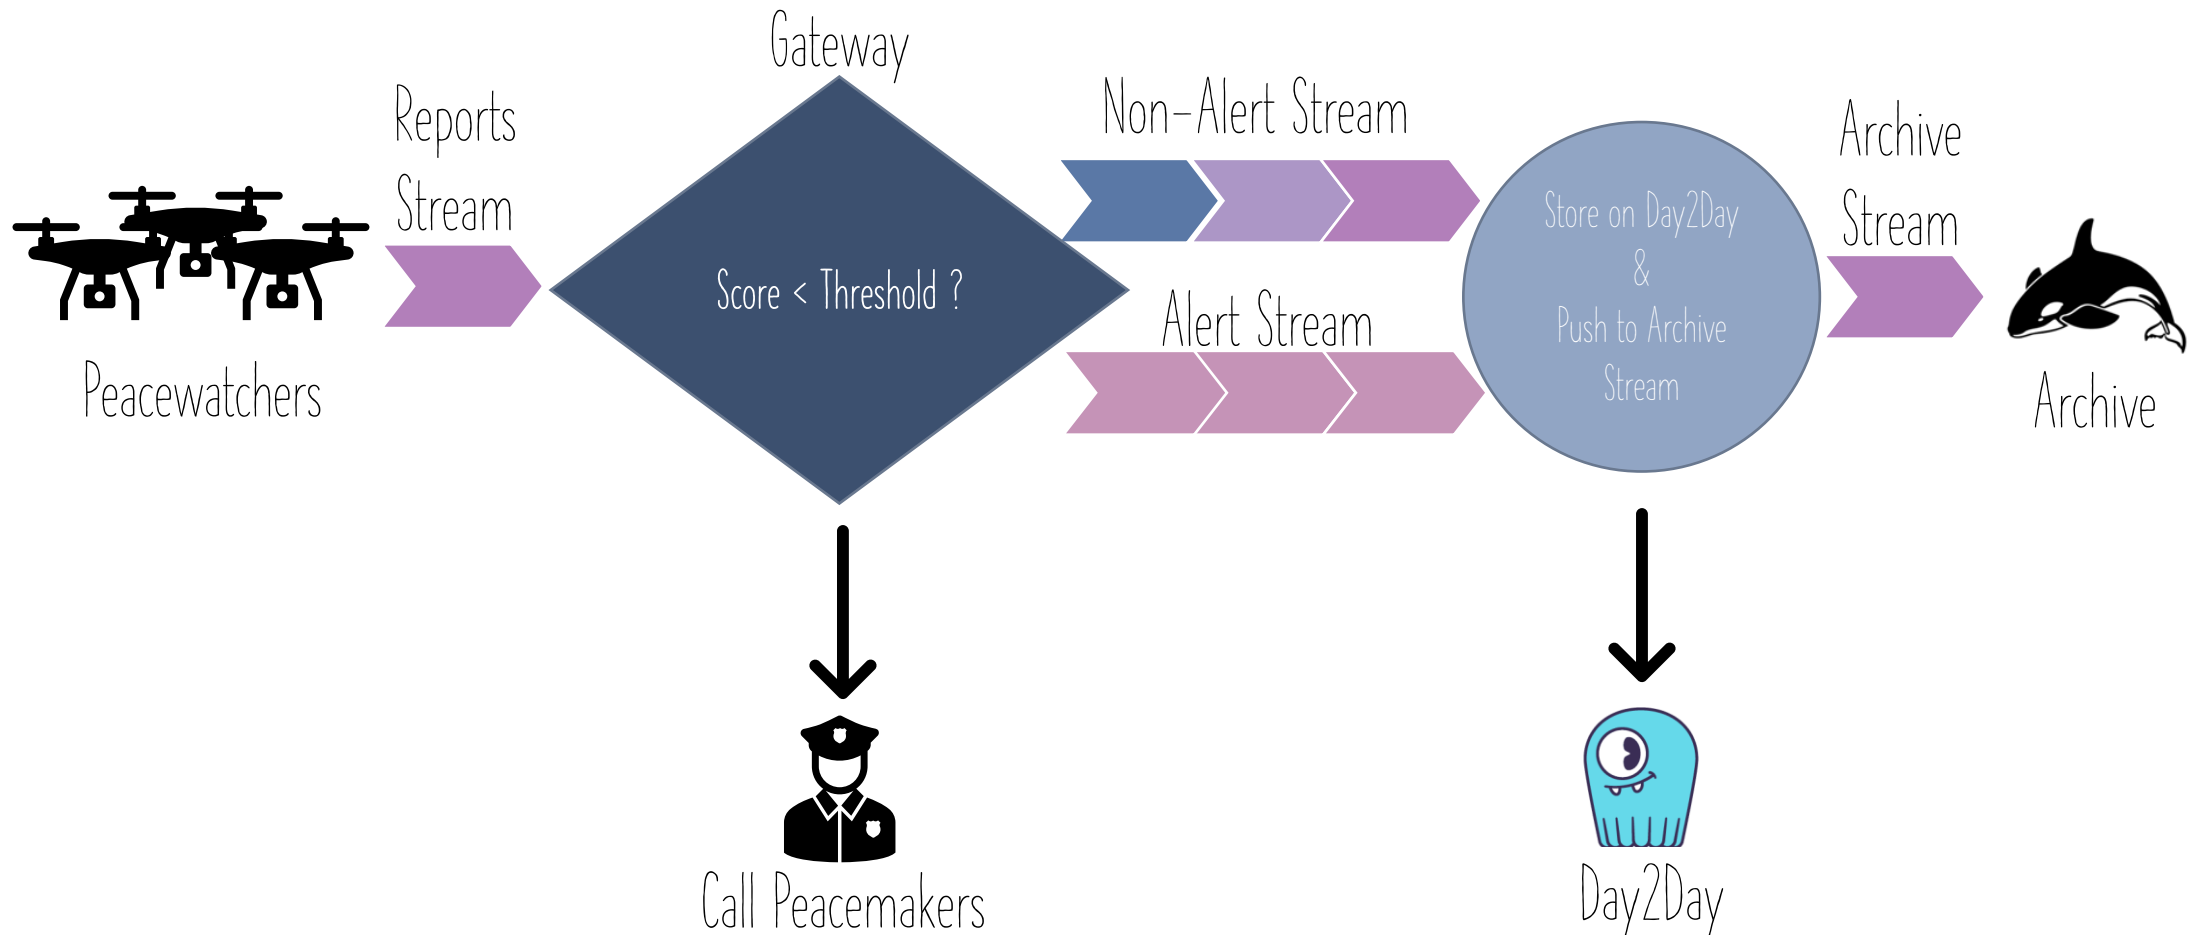
Timestamp of the scan

Unique Identifiers for citizens *(to avoid collisions)*

False positive handling *(can we call peacemakers on a false positive?)*

Peace score delta *(score of the action rather than that of the person)*

OUR SOLUTION'S ARCHITECTURE



ANY QUESTIONS ?

